

gesis

Leibniz Institute
for the Social Sciences



Automatic Sampling and Analysis of YouTube Data

Basic text analysis of user comments

Johannes Breuer, Annika Deubel, & M. Rohangis Mohseni

February 15th, 2023

Required Libraries for This Session

```
library(tidyverse)
library(lubridate)
library(tuber)
library(quanteda)
library(quanteda.textstats)
library(wordcloud)
```

We also need two libraries that are only available from *GitHub*. You can install them using the `install_github()` function from the `remotes` package.

```
library(remotes)
install_github("dill/emoGG")
install_github("hadley/emo")
library(emoGG)
library(emo)
```

Note: Emil Hvitfeldt has created the `emoji` package which is based on the `emo` package and also available via *CRAN*.

Get the Data

As in the last session, we will be working with the - now processed and cleaned - comments for the **Emoji Movie Trailer**. In case you have collected and saved the comments before, you can just load them at this point.

```
FormattedComments <- readRDS("./data/ParsedEmojiComments.rds")
```

Note: Depending on where you saved the data, how you named the file, and what your current working directory is, you might have to adjust the file path.

Repetition: Collecting Data

If you have not collected and parsed the comments before, you, of course, need to do that before you can analyse any data.

NB: To save time and your *YouTube* API quota limit you might not want to do this now.

Step 1: Collecting the comments

```
Comments <- get_all_comments(video_id="r8pJt4dK_s4") # takes a while
```

Repetition: Parsing the Comments

To run the following code the script `yt_parse.R` as well as the ones containing the helper functions (`CamelCase.R`, `ExtractEmoji.R`, and `ReplaceEmoji.R`) need to be in the working directory (you can find those files in the `scripts` folder in the workshop materials).

```
source("yt_parse.R")  
FormattedComments <- yt_parse(Comments) # this will take a while
```

Note: As an alternative to sourcing the `yt_parse.R` file you could also "manually" run the code from the slides for the session on *Processing and Cleaning User Comments* on the collected comments.

Comments Over Time: Data Wrangling



For a first exploratory plot, we want to plot the development of the number of comments per week over time and show until when 50%, 75%, 90%, and 99% of the comments had been posted. This requires some data wrangling.

```
FormattedComments <- FormattedComments %>%
  arrange(Published) %>%
  mutate(date = date(Published),
         week = floor_date(date,
                           unit = "week",
                           week_start = getOption("lubridate.week.start", 1)),
         counter = 1)

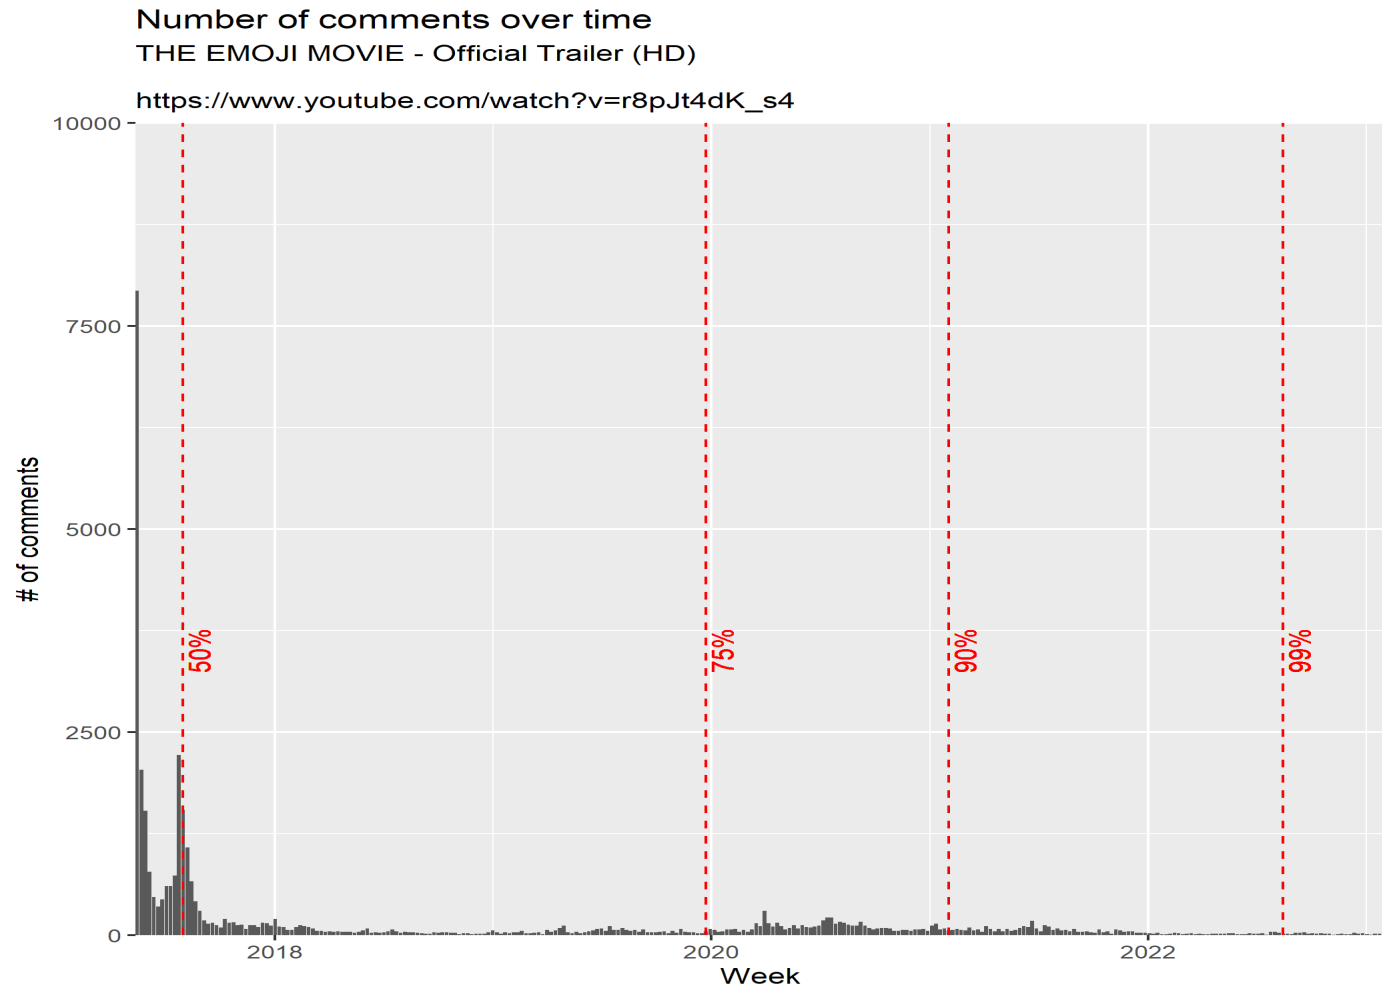
weekly_comments <- FormattedComments %>%
  count(week) %>%
  mutate(cumulative_count = cumsum(n))

PercTimes <- round(quantile(cumsum(FormattedComments$counter),
                           probs = c(0.5, 0.75, 0.9, 0.99)))
```

Comments Over Time: Plot

```
weekly_comments %>%
  ggplot(aes(x = week, y = n)) +
  geom_bar(stat = "identity") +
  scale_x_date(expand = c(0,0)) +
  scale_y_continuous(expand = c(0,0),
                     limits = c(0,10000)) +
  labs(title = "Number of comments over time",
       subtitle = "THE EMOJI MOVIE - Official Trailer (HD)
       \nhttps://www.youtube.com/watch?v=r8pJt4dK\_s4",
       x = "Week",
       y = "# of comments") +
  geom_vline(xintercept = FormattedComments$week[PercTimes], linetype = "dashed",
  geom_text(aes(x = FormattedComments$week[PercTimes][1], label = "50%", y = 3500,
    colour="red", angle=90, vjust = 1.2) +
  geom_text(aes(x = FormattedComments$week[PercTimes][2], label = "75%", y = 3500,
    colour="red", angle=90, vjust = 1.2) +
  geom_text(aes(x = FormattedComments$week[PercTimes][3], label = "90%", y = 3500,
    colour="red", angle=90, vjust = 1.2) +
  geom_text(aes(x = FormattedComments$week[PercTimes][4], label = "99%", y = 3500,
    colour="red", angle=90, vjust = 1.2)
```

Number of Comments Over Time: Plot



Most Popular Comments

Which comments received the highest number of likes?

```
FormattedComments %>%
  arrange(-LikeCount) %>%
  head(10) %>%
  select(Text, LikeCount, Published)
```

	Text	LikeCount	Published
## 1	Will they show Snapchat nudes in the movie?	4287	2017-05-16 15:38:40
## 2	The Meme Movie: Coming 2020	3260	2017-10-16 04:08:12
## 3	Lmao the egg plant emoji never gets used? Do your research lmao	2952	2017-05-16 23:55:38
## 4	The book is so much better because it doesn't exist.	2883	2020-10-30 15:08:17
## 5	I believe everyone intentionally looked this up to dislike it	1802	2020-12-23 18:32:29
## 6	This movie reeks of board room meetings on what kids find "cool".	1690	2017-05-16 22:40:13
## 7	The eggplant emoji never used? Suuuuuree.	1399	2017-05-17 03:10:34
## 8	So, this thing is still a thing? Ugh, I can't really still believe that you cancelled that Popeye movie...	1281	2017-05-16 15:32:41
## 9	This is the best part 2:38	1109	2020-06-08 18:29:03
## 10	So apparently, the eggplant emoji is an emoji that never gets used? Shows how in touch the writers of this movie are.	1090	2021-05-02 14:35:05

Text Mining

An introduction to text mining and analysis (for social sciences) is beyond the scope of this workshop, but there are great introductions available (for free) online, e.g.

- [Text Mining with R - A Tidy Approach](#) by Julia Silge & David Robinson: A tidy(verse) approach
- [Tutorials for the package `quanteda`](#)
- [Text mining for humanists and social scientists in R](#) by Andreas Niekler & Gregor Wiedemann
- [Text Mining in R](#) by Jan Kirenz
- [Computational Text Analysis](#) by Theresa Gessler
- [Automated Content Analysis](#) by Chung-hong Chan

In the following, we will very briefly introduce some key terms and steps in text mining, and then go through some examples of exploring *YouTube* comments (text + emojis).

Popular Text Mining Packages

- **tm**: the first comprehensive text mining package for R
- **tidytext**: tidyverse tools & tidy data principles
- **quanteda**: very powerful text mining package with extensive documentation

Text as Data (in a 🍷)

Document = collection of text strings

Corpus = collection of documents (+ metadata about the documents)

Token = part of a text that is a meaningful unit of analysis (often individual words)

Vocabulary = list of all distinct words from a corpus (i.e., all types)

Document-term matrix (DTM) or Document-feature matrix (DFM) = matrix with n = # of documents rows and m = size of vocabulary columns where each cell contains the count of a particular word for a particular document

Preprocessing (in a 🍷)

For our examples in this session, we will go through the following preprocessing steps:

1. Basic string operations:

- Transforming to lower case
- Detecting and removing certain patterns in strings (e.g., punctuation, numbers or URLs)

2. Tokenization: Splitting up strings into words (could also be combinations of multiple words: n-grams)

3. Stopword removal: Stopwords are very frequent words that appear in almost all texts (e.g., "a", "but", "it", "the") but have low informational value for most analyses (at least in the social sciences)

Preprocessing (in a 🍷)

NB:

- There are many other preprocessing options that we will not use for our examples, such as **stemming**, **lemmatization** or natural language processing pipelines (e.g., to detect and select specific word types, such as nouns and adjectives).
- Keep in mind that the choice and order of these preprocessing steps is important and should be informed by your research question.

Tokenization

Before we tokenize the comments, we want to remove newline commands from the strings.

```
FormattedComments <- FormattedComments %>%  
  mutate(TextEmojiDeleted = str_replace_all(TextEmojiDeleted,  
                                             pattern = "\\n",  
                                             replacement = " "))
```

Tokenization

Now we can tokenize the comments and remove punctuation, symbols, numbers, and URLs.

```
toks <- FormattedComments %>%  
  pull(TextEmojiDeleted) %>%  
  char_tolower() %>%  
  tokens(remove_numbers = TRUE,  
          remove_punct = TRUE,  
          remove_separators = TRUE,  
          remove_symbols = TRUE,  
          split_hyphens = TRUE,  
          remove_url = TRUE)
```


Document-Feature Matrix

With the tokens we can create a **document-feature matrix** (DFM) and remove **stopwords**.

```
commentsDfm <- dfm(toks,  
                    remove = quanteda::stopwords("english"))
```

Most Frequent Words

```
TermFreq <- textstat_frequency(commentsDfm)
head(TermFreq, n = 20)
```

##	feature	frequency	rank	docfreq	group
## 1	movie	11706	1	8876	all
## 2	emoji	3191	2	2759	all
## 3	like	2816	3	2427	all
## 4	just	2486	4	2199	all
## 5	nom	2239	5	1	all
## 6	people	1541	6	1332	all
## 7	sony	1508	7	1397	all
## 8	bad	1395	8	1275	all
## 9	good	1312	9	1198	all
## 10	one	1227	10	1111	all
## 11	hate	1127	11	1019	all
## 12	emojis	1099	12	990	all
## 13	see	1047	13	927	all
## 14	watch	1030	14	948	all
## 15	make	1010	15	908	all
## 16	think	992	16	898	all
## 17	know	960	17	879	all
## 18	popeye	939	18	862	all
## 19	dislikes	912	19	895	all
## 20	can	886	20	772	all

Removing Tokens

We may want to remove additional words (that are not included in the stopwords list) if we consider them irrelevant for our analyses.

```
custom_stopwords <- c("nom", "just", "one")
commentsDfm <- dfm(toks, remove = c(quanteda::stopwords("english"),
                                   custom_stopwords))
TermFreq <- textstat_frequency(commentsDfm)
```

For more options for selecting or removing tokens, see the [quanteda documentation](#).

Wordclouds

```
wordcloud(words = TermFreq$feature,
          freq = TermFreq$frequency,
          min.freq = 10,
          max.words = 50,
          random.order = FALSE,
          rot.per = 0.35,
          colors = brewer.pal(8, "Dark2"))
```

Note: You can adjust what is plotted by, e.g., changing the minimum frequency (`min.freq`) or the maximum # of words (`max.words`). Check `?wordcloud` for more customization options.

Wordclouds



Plot Most Frequent Words

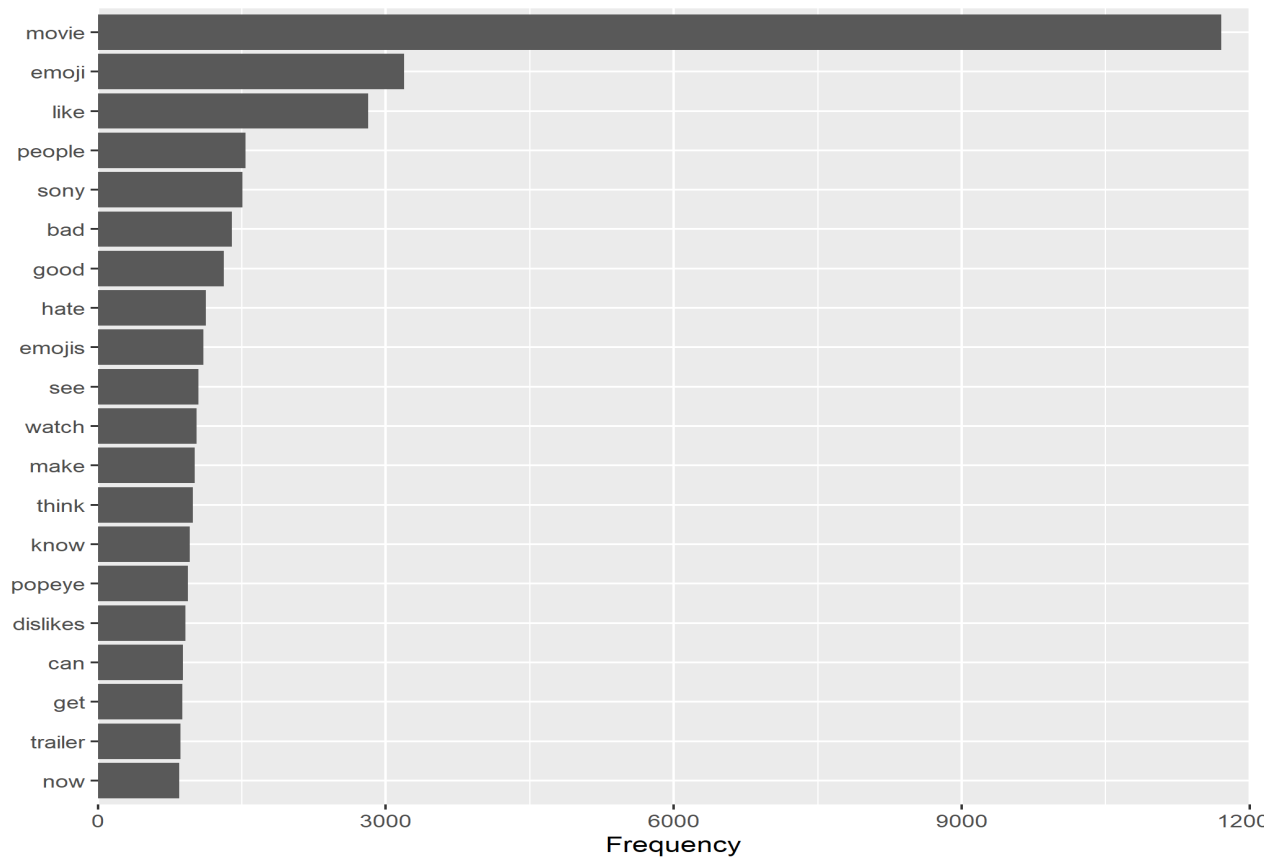
```
TermFreq %>%
head(n = 20) %>%
  ggplot(aes(x = reorder(feature, frequency), y = frequency)) +
  geom_bar(stat="identity") +
  labs(title = "Most frequent words in comments",
        subtitle = "THE EMOJI MOVIE - Official Trailer (HD)
\nhttps://www.youtube.com/watch?v=r8pJt4dK_s4",
        x = "",
        y = "Frequency") +
  scale_y_continuous(expand = c(0,0),
                     limits = c(0,12000)) +
  coord_flip()
```

Plot Most Frequent Words

Most frequent words in comments

THE EMOJI MOVIE - Official Trailer (HD)

https://www.youtube.com/watch?v=r8pJt4dK_s4



Plot Docfreq

Instead of the raw frequency of words we can also look at the number of comments that a particular word appears in. This metric takes into account that words might be used multiple times in the same comment.

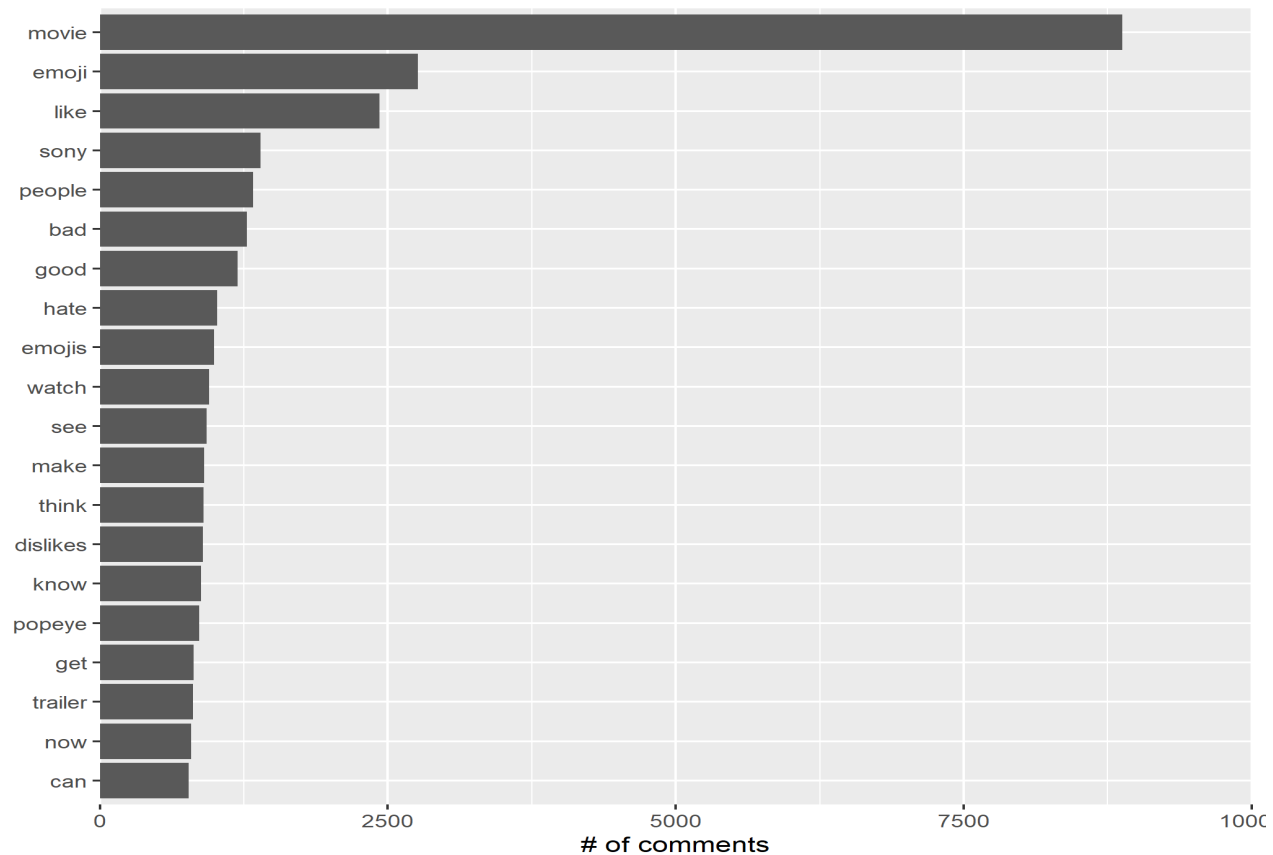
```
TermFreq %>%
head(n = 20) %>%
  ggplot(aes(x = reorder(feature, docfreq), y = docfreq)) +
  geom_bar(stat="identity") +
  labs(title = "Words that appear in the highest number of comments",
        subtitle = "THE EMOJI MOVIE - Official Trailer (HD)
        \nhttps://www.youtube.com/watch?v=r8pJt4dK\_s4",
        x = "",
        y = "# of comments") +
  scale_y_continuous(expand = c(0,0),
                     limits = c(0,10000)) +
  coord_flip()
```


Plot Docfreq

Words that appear in the highest number of comments

THE EMOJI MOVIE - Official Trailer (HD)

https://www.youtube.com/watch?v=r8pJt4dK_s4



Emojis

In most of the research studying user-generated text from social media, emojis have, so far, been largely ignored. However, emojis convey emotions and meaning, and can, thus, provide additional information or context when working with textual data.

In the following, we will do some exploratory analysis of emoji frequencies in *YouTube* comments. Before we can start, we first need to do some data cleaning again, then tokenize the emojis as some comments include more than one emoji, and create an emoji DFM.

```
emoji_toks <- FormattedComments %>%  
  mutate(Emoji = na_if(Emoji, "NA")) %>% # define missings  
  mutate (Emoji = str_trim(Emoji)) %>% # remove spaces  
  filter(!is.na(Emoji)) %>% # only keep comments with emojis  
  pull(Emoji) %>% # pull out column cotaining emoji labels  
  tokens(what = "fastestword") # tokenize emoji labels  
  
EmojiDfm <- dfm(emoji_toks) # create DFM for emojis
```

Most Frequent Emojis

```
EmojiFreq <- textstat_frequency(EmojiDfm)
head(EmojiFreq, n = 10)
```


##		feature	frequency	rank	docfreq	group
## 1		emoji_pileofpoo	3999	1	536	all
## 2		emoji_eggplant	3578	2	279	all
## 3		emoji_facewithtearsofjoy	2873	3	853	all
## 4		emoji_unamusedface	2488	4	677	all
## 5		emoji_bbutton(bloodtype)	1878	5	130	all
## 6		emoji_middlefinger	1866	6	298	all
## 7		emoji_grinningface	1562	7	374	all
## 8		emoji_flushedface	1239	8	256	all
## 9		emoji_thumbsdown	1157	9	264	all
## 10		emoji_facewithsymbolsonmouth	982	10	104	all

Plot Most Frequent Emojis

```
EmojiFreq %>%
head(n = 10) %>%
  ggplot(aes(x = reorder(feature, frequency), y = frequency)) +
  geom_bar(stat="identity") +
  labs(title = "Most frequent emojis in comments",
        subtitle = "THE EMOJI MOVIE - Official Trailer (HD)
\nhttps://www.youtube.com/watch?v=r8pJt4dK_s4",
        x = "",
        y = "Frequency") +
  scale_y_continuous(expand = c(0,0),
                      limits = c(0,5000)) +
  coord_flip()
```

Note: Similar to what we did for the comment text before we could replace `frequency` with `docfreq` in the above code to create a plot with the emojis that appear in the highest number of comments.

Plot Most Frequent Emojis



Emoji Frequency Plot: Preparation (1)

The previous emoji frequency plot was a bit 🥲. To make things prettier, we can use the actual emojis instead of the text labels in our plot. Doing this takes a bit of preparation.¹

As a first step, we need an emoji lookup table in which the values in the name column have the same format as the labels in the feature column of our `EmojiFreq` object.

```
emoji_lookup <- jis %>%
  select(runes, name) %>%
  mutate(runes = str_to_lower(runes),
         name = str_to_lower(name)) %>%
  mutate(name = str_replace_all(name, " ", "")) %>%
  mutate(name = paste0("emoji_", name))
```

¹For an alternative approach to using emojis in `ggplot2` see this [blog post by Emil Hvitfeldt](#).

Emoji Frequency Plot: Preparation (2)

The second step of preparation for the nicer emoji frequency plot is creating mappings of emojis to data points so that we can use emojis instead of points in a scatter plot.¹

```
top_emojis <- 1:10

for(i in top_emojis){
  name <- paste0("mapping", i)
  assign(name,
    do.call(ggeom_emoji, list(data = EmojiFreq[i,],
                             emoji = gsub("^0{2}", "", strsplit(tolower(emoji_
```

¹ Please note: this code has not been tested systematically. Depending on which emojis are most frequent for a video, this might not work because (a) one of the emojis is not included in the emoji lookup table (which uses the `jis` data frame from the `emo` package) or (b) the content in the `runes` column doesn't match the format/code that the `emoji` argument in the `ggeom_emoji` function from the `emoGG`



Emoji Frequency Plot

```
EmojiFreq %>%
head(n = 10) %>%
  ggplot(aes(x = reorder(feature, -frequency), y = frequency)) +
  geom_bar(stat="identity",
           color = "black",
           fill = "#FF74A6",
           alpha = 0.7) +
  geom_point() +
  labs(title = "Most frequent emojis in comments",
       subtitle = "THE EMOJI MOVIE - Official Trailer (HD)
\nhttps://www.youtube.com/watch?v=r8pJt4dK_s4",
       x = "",
       y = "Frequency") +
  scale_y_continuous(expand = c(0,0),
                     limits = c(0,5000)) +
  theme(panel.grid.major.x = element_blank(),
        axis.text.x = element_blank(),
        axis.ticks.x = element_blank()) +
  mapping1 +
  mapping2 +
  mapping3 +
  mapping4 +
  mapping5 +
  mapping6 +
  mapping7 +
  mapping8 +
  mapping9 +
  mapping10
```

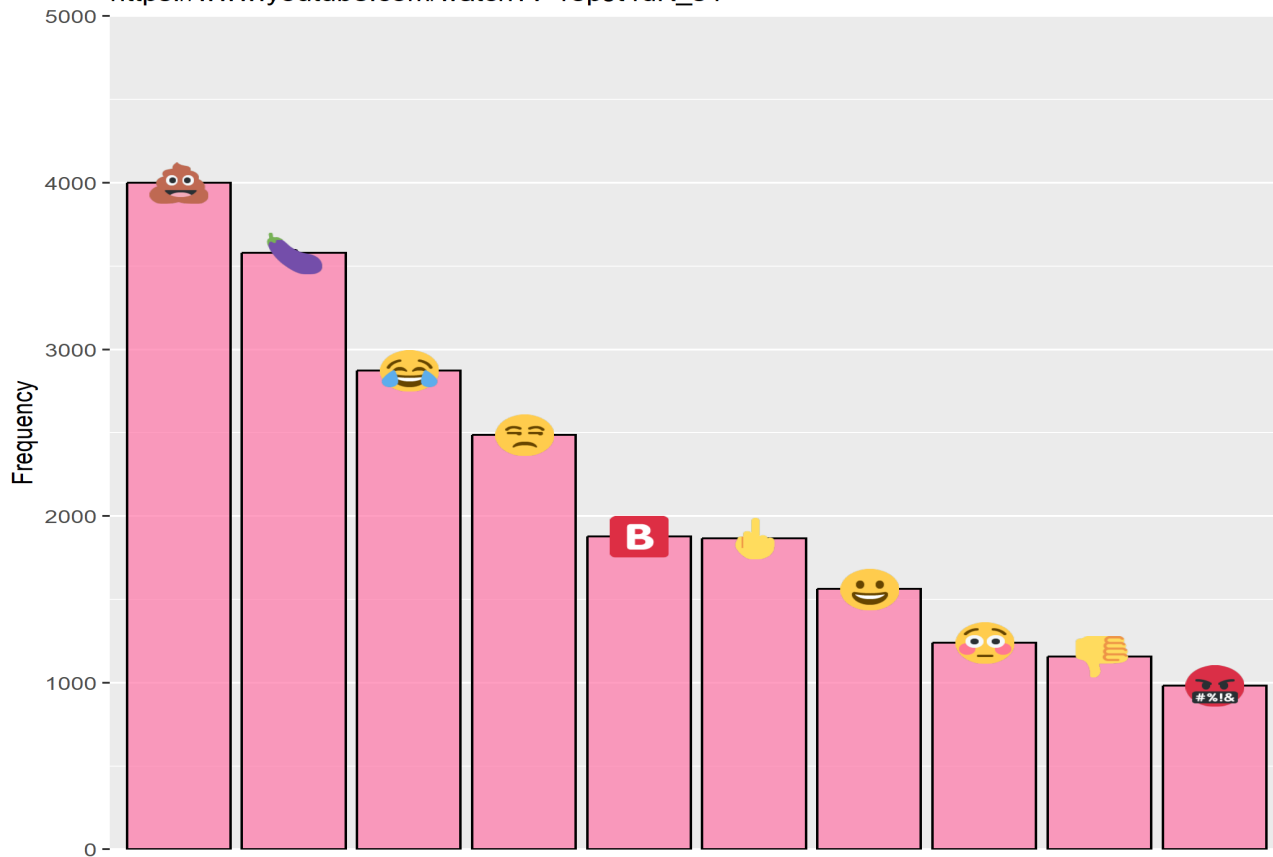



Emoji Frequency Plot

Most frequent emojis in comments

THE EMOJI MOVIE - Official Trailer (HD)

https://www.youtube.com/watch?v=r8pJt4dK_s4



Exercise time 🏋️ 💪 🏃 🚴

Solutions